

Dynamic Image Generation For Action Recognition

...

Team Member

Khushi Agarwal

•••
Under the Supervision of
Dr. Aruna Tiwari Professor,
CSE

Problem Statement

Videos outnumber still images by a huge margin in terms of the amount of visual data that is now available. Therefore, it is crucial to precisely and thoroughly comprehend the content of videos. Videos are made up of a series of still pictures. The challenge of comprehending video content is the main concern. This is accomplished by viewing films as a stack of frames, after which discriminative models are developed that extract the data required to address certain issues like action recognition.

Continued.....

In this project, I have provided a new method for representing videos that, like the preceding instances, encodes the data in a generic and content-independent way, producing a long-lasting, reliable motion representation useful for both action recognition and other video analysis tasks. This fresh depiction simplifies a single image, which we refer to as the dynamic image, that combines the motion information present in each frame of a movie.

What is Dynamic Image ??

- Dynamic image generation from video refers to the process of extracting still images or frames from a video and then generating new images based on those frames.
- It is a RGB still image that summarizes, in a compressed format, the gist of a whole video (sub)sequence
- They explain some form of movement.



What exactly is being done in our project

We need to develop a dynamic image for a video that will be utilised for action recognition later on. We will build dynamic images from a video by combining a certain number of frames. For example, if we have a 25-frame video, we will generate 1 dynamic image using these frames. Now, using that dynamic image, the video will then be classified based on human activity, i.e. action recognition. (This will be handled in the testing part).

How Frames Combined to generate a single image ?

We have taken the mean of all the frames for a video.

We also tried some other methods but the mean was giving good results hence proceeded with that.

Data Set


We have used UCF101 (subset of its classes) datasets.

- In UCF-101 we have 13, 320 clips distributed in 101 action classes.
- For our project we have taken its 9 classes for the classification purpose.

-

- ApplyEyeMakeup
 - ApplyLipstick
 - Archery
 - BabyCrawling
 - BalanceBeam
 - BandMarching
 - BaseballPitch
 - Basketball
 - BasketballDunk

9 classes with 1214 images in total



Directory Structure of UCF-101

UCF-101 (101 directories)



ApplyEyeMakeup
145 files



ApplyLipstick
114 files



Archery
145 files



BabyCrawling
132 files



BalanceBeam
108 files



BandMarching
155 files



BaseballPitch
150 files



Basketball
134 files



BasketballDunk
131 files



BenchPress
160 files



Biking
134 files



Billiards
150 files

Example of video from the dataset



Data Preprocessing

Data preprocessing is the process of cleaning, transforming, and preparing raw data into a format that is suitable for analysis. It is a crucial step in any data analysis project, as the quality of the data input directly affects the quality of the output of any analysis.

The specific techniques used in data preprocessing depend on the nature of the data, the intended use of the data, and the analysis methods to be employed. Data preprocessing is often considered one of the most time-consuming and resource-intensive steps in a data analysis project, but it is also one of the most critical steps to ensure accurate and reliable results.

Steps

Data preprocessing involves several steps, including:

- **Data cleaning:** Removing or correcting any errors, inconsistencies, or missing values in the data.
- **Data transformation:** Converting data into a more suitable format, such as normalizing or scaling data, converting categorical data into numerical data, or reducing the dimensionality of the data.
- **Data integration:** Combining multiple data sources into a single dataset.
- **Data reduction:** Reducing the size of the data by selecting a subset of the most relevant features or examples.
- **Data discretization:** Converting continuous data into discrete categories.

We need to convert video clips into frames

Some basic steps that we are going to implement for obtaining frames :

1. **Load the video clip:** The first step is to load the video clip using a suitable library such as OpenCV or FFmpeg. These libraries provide various functions for working with video files.
2. **Extract the frames:** Once the video clip is loaded, the next step is to extract the frames from the clip. This can be done by looping through the video frames and saving each frame as an image file.
3. **Preprocess the frames:** After extracting the frames, you may need to preprocess them to improve the quality of the images. This could include tasks such as resizing, cropping, and color correction.
4. **Save the frames:** Once the frames have been extracted and preprocessed, they need to be saved in a suitable format such as JPEG or PNG.
5. **Use the frames for analysis:** Once the frames have been saved, you can use them for various analysis tasks such as object detection, image classification, or image segmentation.

Preprocessing part

1. Used Open CV libraries video capture function to extract the frames from the video at certain intervals.
2. The image read by cv2 was in BGR format, we have to convert it into RGB format. Without this option images will be shown in blue hue because matplotlib uses RGB to display image
3. Resizing all the images to the same dimensions (330 x 330).
4. Cropping edges of images using Canny algorithm.
5. Now the images are ready for training and testing.

Some of the dynamic images generated with their labels on the top

BandMarching: 6



Archery: 5




BaseballPitch: 2



BalanceBeam: 4



ApplyEyeMakeup: 3



Archery: 5



BaseballPitch: 2



Basketball: 8



BandMarching: 6




BaseballPitch: 2



BabyCrawling: 7



ApplyEyeMakeup: 3



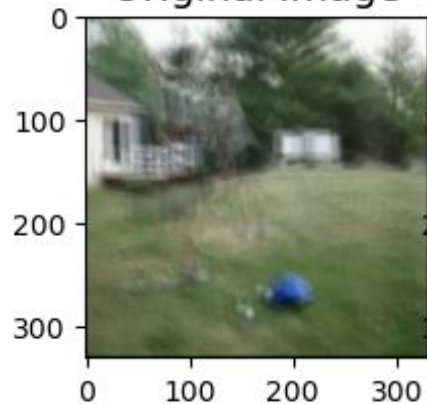
Canny algorithm

The Canny algorithm is a way to find edges in a digital image. Edges are the places where the image changes rapidly from one color to another, like the edges of a book or a person's face. The algorithm works by following these steps


1. It smooths out the image to remove any small bumps or jitters that might look like edges, so it can focus on the real edges.
2. It looks for places where the colors in the image change a lot, and figures out which direction the change is happening in.
3. It focuses on the strongest edges, ignoring any weaker edges that might be caused by noise or other factors.
4. Finally, it links up any weak edges that are close to strong edges, to create a complete picture of all the edges in the image.

The output of the algorithm is a picture where the edges are shown as lines or curves, and everything else is black. The Canny algorithm is used in many different applications, like recognizing objects in pictures, tracking motion, etc.

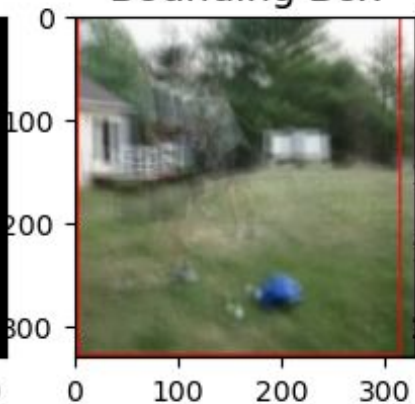
Original Image




Canny Edges




Bounding Box




Cropped




Original Image




Canny Edges



Bounding Box




Cropped



Data Augmentation

Image augmentation is a creation of additional training data based on existing images, for example translation, rotation, flips and zoom. Using ImageDataGenerator class from Keras library create additional images of each class.



Convolutional Neural Network

CNNs are fully connected feed forward neural networks.

CNNs are trained to identify the parameters like edges of objects in any image.

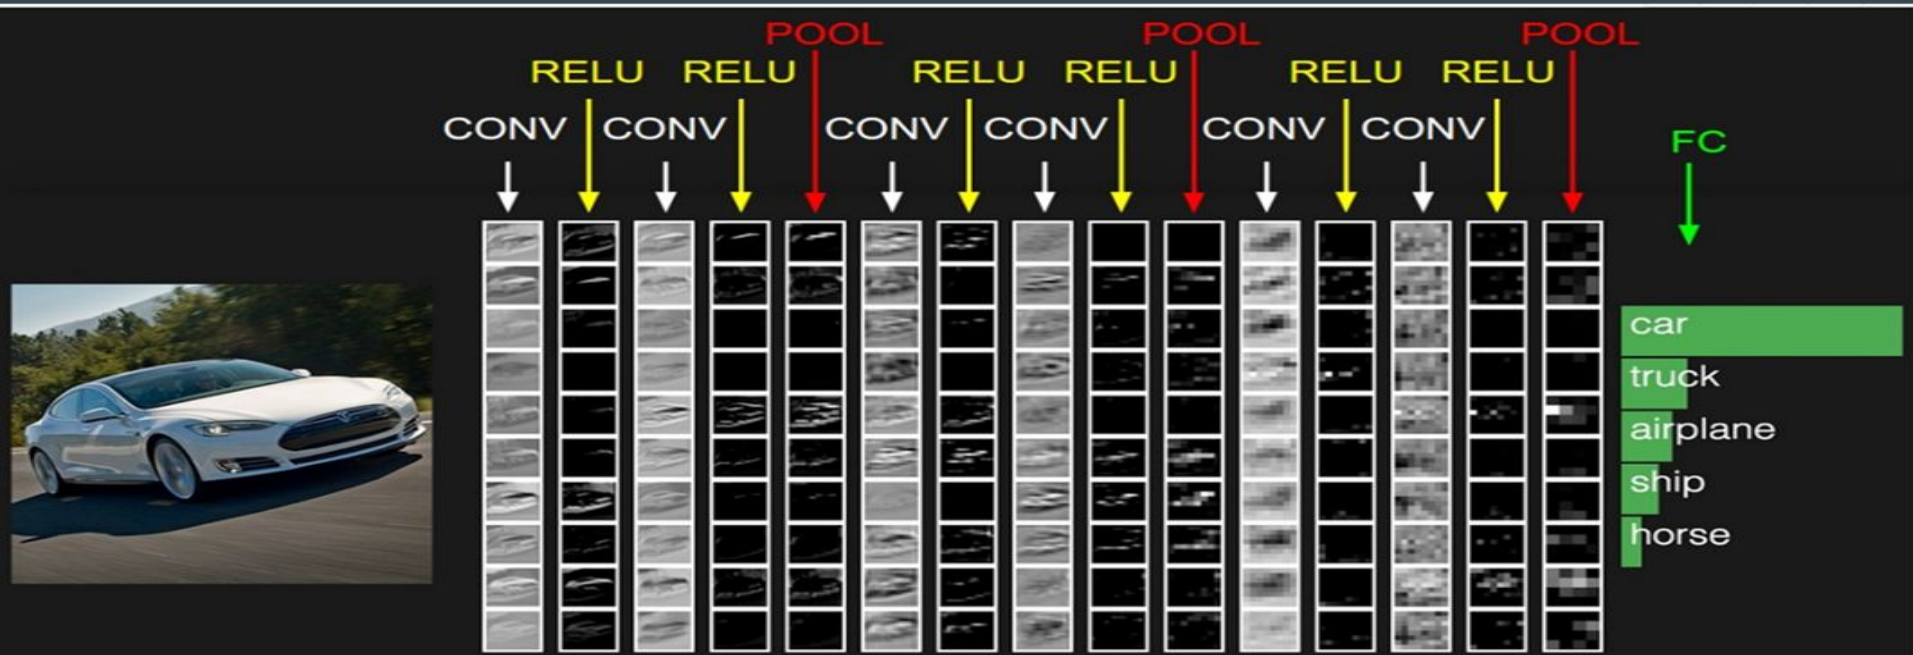
CNNs are very effective in reducing the number of parameters without losing on the quality of models. Images have high dimensionality (as each pixel is considered as a feature) which suits the above described abilities of CNNs

CNN is mainly used to extract features from the images.

It contains a set of kernels (filters), parameters which are to be learned throughout the trainings.

Each filter convolves with the image and feature maps are created.

A simple CNN structure



CONV: Convolutional kernel layer
RELU: Activation function
POOL: Dimension reduction layer
FC: Fully connection layer

Architecture of the CNN Model used

Layers

- 1 Conv2D 32 -> Pool
- 2 Conv2D 64 -> Pool
- 3 Conv2D 128 -> Pool
- 4 Conv2D 128 -> Pool
- 5 Conv2D 128 -> Pool
- 6 FLAT
- 7 Drop
- 8 Dense 512
- 9 Dense len(CLASSES)

Continued.....

- Conv2D with hyperparameters mentioned above: Conv2D(kernel_size, (filters, filters), input_shape=(img_w, img_h, 3)) with activation function for each layer as a Rectified Linear Unit (ReLU): Activation('relu').
- MaxPooling2D layer to reduce the spatial size of the incoming features; 2D input space: MaxPooling2D(pool_size=(max_pool, max_pool)).
- Do the same increasing the kernel size: 32 -> 64 -> 128 -> 128 -> 128
- Flatten the input: transform the multidimensional vector into a single dimensional vector: Flatten().
- Add dropout layer which randomly sets a certain fraction of its input to 0 and helps to reduce overfitting: Dropout(0.5)

Continued.....

- Add fully connected layer with 512 nodes and activation function relu: `Dense(512), Activation('relu')`.
- Provide last fully connected layer which specifies the number of classes . Softmax activation function outputs a vector that represents the probability distributions of a list of potential outcomes: `Dense(activation='softmax')`.

Model Summary

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 220, 220, 32)	896
max_pooling2d (MaxPooling2D)	(None, 110, 110, 32)	0
conv2d_1 (Conv2D)	(None, 110, 110, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 55, 55, 64)	0
conv2d_2 (Conv2D)	(None, 55, 55, 128)	73856
max_pooling2d_2 (MaxPooling2D)	(None, 27, 27, 128)	0
conv2d_3 (Conv2D)	(None, 27, 27, 128)	147584
average_pooling2d (AveragePooling2D)	(None, 13, 13, 128)	0
conv2d_4 (Conv2D)	(None, 13, 13, 128)	147584
max_pooling2d_3 (MaxPooling2D)	(None, 6, 6, 128)	0
flatten (Flatten)	(None, 4608)	0
dropout (Dropout)	(None, 4608)	0
dense (Dense)	(None, 512)	2359808
dense_1 (Dense)	(None, 9)	4617

=====
Total params: 2,752,841
Trainable params: 2,752,841
Non-trainable params: 0

Optimizer used

- ❖ Optimizer that implements the Adam algorithm.
- ❖ Adam optimization is a stochastic gradient descent method that is based on adaptive estimation of first-order and second-order moments.
- ❖ Some of its arguments are :
 - learning_rate : default to 0.001
 - epsilon : defaults to $1e-7$
 - beta_1: the exponential decay rate for the 1st moment estimates, defaults to 0.999.
 - beta_2: the exponential decay rate for the 2nd moment estimates, defaults to 0.999.
 - The first moment is mean, and the second moment is variance.

Continued.....

Overall, this CNN has 5 convolutional layers with increasing filter sizes and 4 max pooling layers to reduce the spatial dimensions of the features. The final dense layers are used for classification. The Dropout layer is added to prevent overfitting.

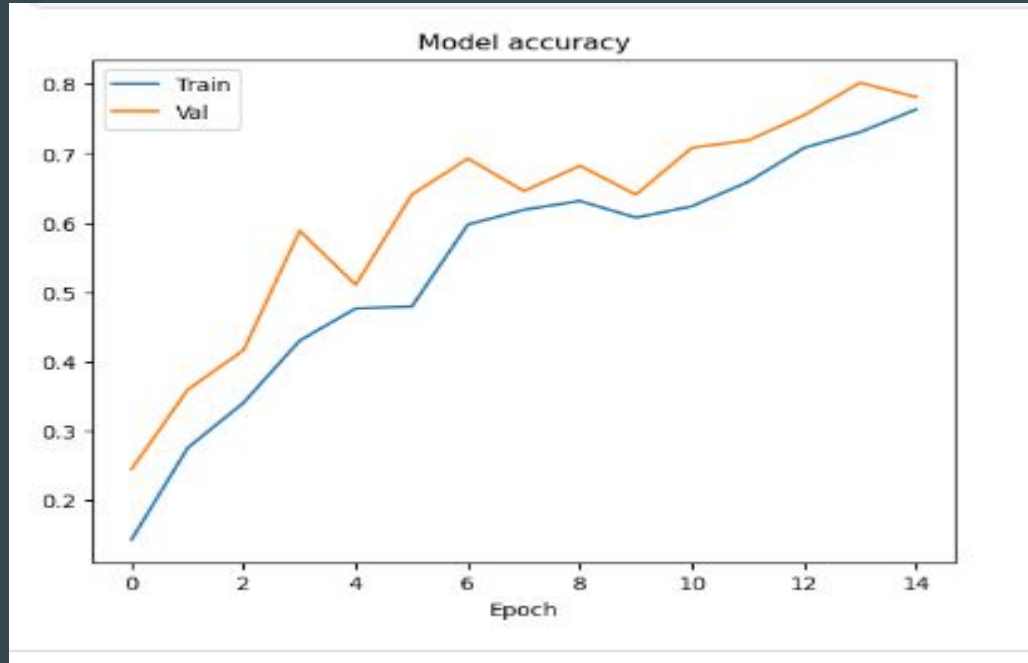
Loss function

- It measures how different the predicted outputs of a machine learning model are from the true outputs. The goal of a deep learning model is to minimize the value of the loss function, as this indicates that the model is making more accurate predictions.
- Here used the loss function “`sparse_categorical_crossentropy`”.
- Sparse categorical cross-entropy is a loss function used in multiclass classification problems where the target variable is represented as integer labels. It measures the difference between the predicted probability distribution over the classes and the actual integer labels. The loss is calculated by taking the negative log of the predicted probability for the correct class label. During training, the model updates its weights in the direction that minimizes the sparse categorical cross-entropy loss function.

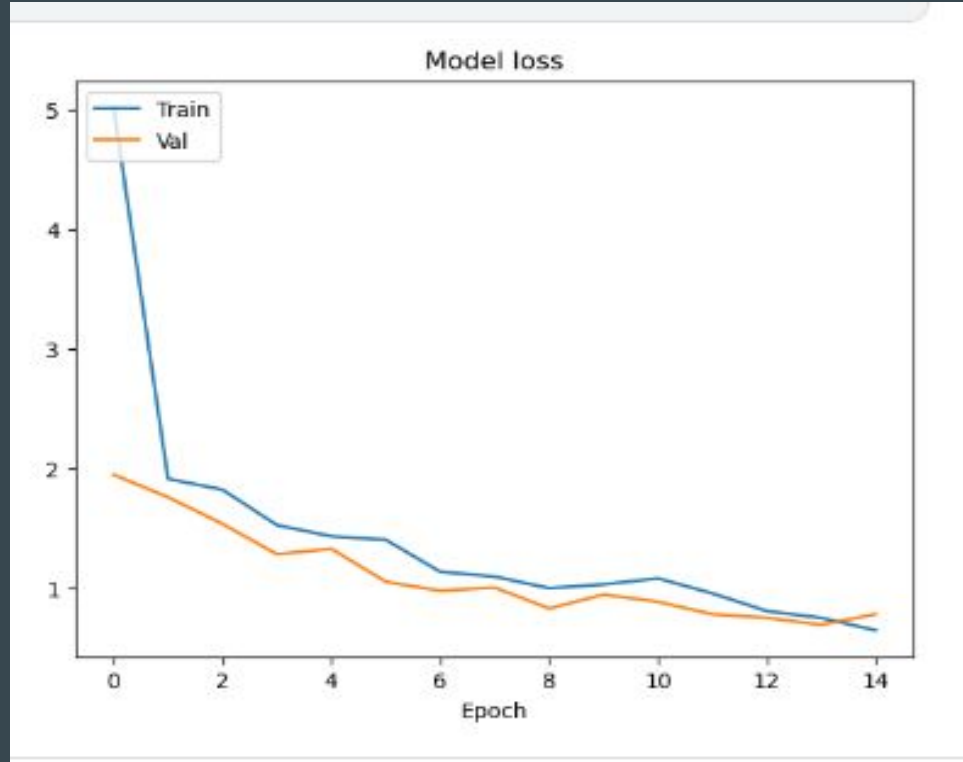
Performance Metrics

- **Accuracy:** The percentage of correctly classified samples out of the total number of samples.
- **Confusion matrix:** A table that summarizes the classification results. It shows the number of true positives, false positives, true negatives, and false negatives.

Model Accuracy on training and validation dataset



Model loss on training and validation dataset



Loss and Accuracy on training dataset

After 15 epochs.....


- loss: 0.6511
- accuracy: 0.7635
- val_loss: 0.7868
- val_accuracy: 0.7812

Loss and Accuracy on validation dataset.....

```
loss:0.774970293045044
```

```
accuracy:0.7777777910232544
```


Confusion Matrix For Validation Data Results




Test data Accuracy

Accuracy : 78.33333333333333

Confusion Matrix For Test Data Results



Some of the images from the test folder with a label, class which model predicted, and actual class.



Thank you !!